

Real or Not? NLP with Disaster Tweets EDA

Vegar Andreas Bergum, *vab1g19*, ChanKai Fu, *ckf1n19*, Adam Ghoumrassi, *ag3u19*, and PingChun Tsai, *pct1g19*.

Abstract—Lorem ipsum dolor sit amet, consectetur adipiscing elit. Fusce eget erat mi. Nam non diam felis. Ut tincidunt, mauris eget ornare sodales, orci ligula bibendum dui, id tristique augue arcu eu dolor. Vestibulum tortor justo, malesuada at ligula eget, cursus laoreet diam. Orci varius natoque penatibus et magnis dis parturient montes, nascetur ridiculus mus. Ut non volutpat metus, id luctus nunc. Duis nisi nunc, ullamcorper at enim at, maximus luctus sem. Quisque ipsum risus, tempor nec sem vitae, placerat placerat tellus. Nunc congue, diam facilisis eleifend elementum, sapien sem viverra risus, lobortis tincidunt libero leo quis est. Sed porta tellus a egestas fermentum. Suspendisse eget tristique libero, sit amet maximus ligula. Suspendisse in tellus sagittis, ultrices ante quis, tristique lectus. Quisque fringilla, leo sagittis iaculis finibus, nunc nulla elementum neque, quis commodo justo libero a nunc. Integer convallis sapien eu laoreet congue.

I. INTRODUCTION

THIS project considers the *Real or Not? NLP with Disaster Tweets*¹ Kaggle competition, where one attempts to classify whether or not a tweet is announcing a real disaster. As stated in the competition; "Twitter has become an important communication channel in times of emergency", both by government organisations and individuals. With this comes the challenge of fake news and the general authenticity of the information being spread online. With the capability of quickly and reliably classifying any given tweet as announcing a real disaster or not, one can help improve automatic disaster monitoring.

With any machine learning project, one of the most important steps is exploring the data available. In this competition, we have access to 7,613 labelled tweets (fake and real). Most of this data is pure text (the content of the tweets), which makes this a natural language processing (NLP) task. Feature extraction is very important and complex when working with textual natural language. There are several methods for representing the semantics of text, including word-frequency based approaches such as TF-IDF and more advanced latent space representations such as word embeddings like *word2vec* [1]. In order to get a better understanding of the existing features in the data and basis for choosing a semantic representation model, we explore both token- and character-level relationships between the data and the two classes.

II. DATA SUMMARIZATION

The data set contains 7,613 tweets, labelled as fake and real. Excluding the label, the data has three main attributes: *keyword*, *location*, *text*. The keyword attribute contains a list of pre-defined keywords that describes the contents of the tweet.

The location is provided directly by Twitter's geo-tag on a tweet and is represented as a string, e.g. *New York City*. The text attribute holds the vast majority of the data, as this is the raw textual content of the tweet. Table I shows the percentage of missing values and unique entries in each attribute.

Attribute	Missing values	Unique entries
keyword	0.8%	61
location	33.3%	2,533
text	0.0%	7,613

TABLE I: Missing and unique values in the data attributes.

As can be seen from the table, a third of the location data is missing. The keyword attribute is only missing a relatively few number of entries, although it should be noted that the number of unique keywords is quite low. Further, Figure 1 shows the proportions of tweets in both classes. There is a slight unbalance in the dataset, however, this is considered low enough to not have any substantial effects on the problem modeling.




Fig. 1: Class proportions.

III. QUANTITATIVE STATISTICS

Exploring the textual data in the corpus reveals some correlations between the various syntactic and semantic features, and the two classes. Figure 2 shows that there is a small shift in the mean of the average number of words per tweet between the two classes. Non-disaster (fake) tweets are, on average, shorter and have a larger variation in tweet size. This implies that the tweet length could be a useful feature for classification.

Emojis have a prominent presence in social media, also on Twitter. Unfortunately, the text in the dataset is encoded in

¹<https://www.kaggle.com/c/nlp-getting-started/overview>




Fig. 2: Distribution of average number of words per tweet.

such a way that emojis can not be distinguished from each other. However, one can still look at the frequency of their use. *A priori* one could argue that a real disaster tweet is often from a mainstream news source and would stray away from using informal language and symbols, such as emojis. Figure 3, showing the frequency of emoji-usage, confirms this hypothesis. Non-disaster tweets use emojis more regularly than real disaster tweets. Since emojis are encoded as a special escape-sequence they can be considered as a semantic token equivalently to a regular word when modeling the problem for classification.




Fig. 3: Frequency of emoji-usage in the two classes.

The use of external links has shown to be a large differentiator between the two classes. 50% of non-disaster tweets contain external links, compared to 78% of disaster tweets. A rationale for this difference is that disaster tweets often refer to a source or proof of the disaster. Looking into the links being used, we can observe a difference in pages that are linked to from the two classes. Figure 4 shows the top 15 referred-to domains by the two classes. Observe that the disaster tweets are linking to official news sites such as bbc.co.uk, latimes.com, cnn.com, while the non-disaster tweets link

to other social-medias sites such as facebook.com and instagram.com, ebay.com and various blogs.

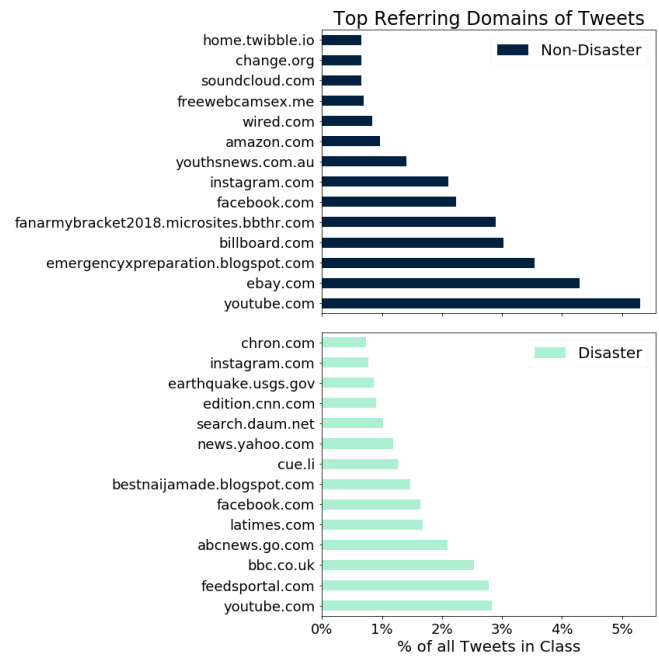


Fig. 4: Top domains for links in tweets.

Punctuation ...




Fig. 5: Use of punctuation per tweet and differences in use between classes.

Figure 6 illustrates some of the different linguistic features of the two classes by showing the distribution and difference in part-of-speech tags per tweet. Notably, the disaster tweets more often use numbers and proper nouns. One explanation of this difference would be the larger amount of references to statistics (numbers) and official bodies and names (proper nouns) in real disaster coverage. However, the difference is very minimal for all POS-tags, so this could simply be an arbitrary difference due to noise. The `en_core_web_sm` POS-tagger in `spacy` was used for tagging.




Fig. 6: Distribution of POS tags per tweet and differences between classes.

Using a Hamilton distance of 1 between two tokens and the WordFrequency project² dictionary to correct spelling, we evaluate the number of spelling mistakes per tweet, per class. Figure 7 shows the number of spelling mistakes for both classes. We observe that except for tweets with a single spelling mistake, disaster tweets contain more mistakes than non-disaster tweets. These differences are very small and lay within the margin of error to be expected from approach to detecting spelling mistakes used.

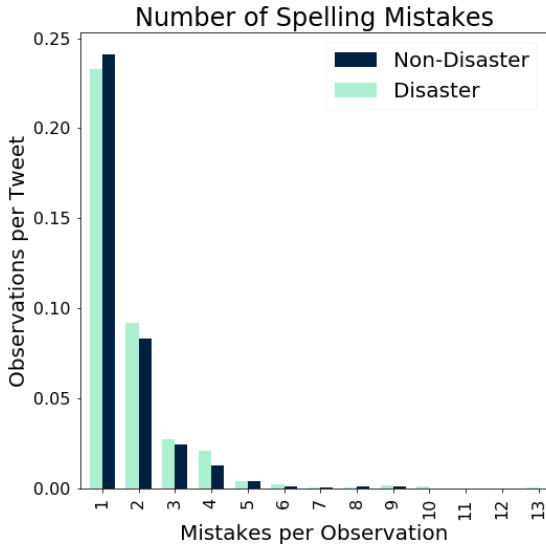


Fig. 7: Frequency of spelling mistakes per tweet.

IV. SENTIMENT ANALYSIS

As we have seen from the investigation into the use of external links, many of the real disaster tweets are referring

to large mainstream news sources. News is supposed to be objective and carry little to no sentiment. Using VADER Sentiment Analysis³, a rule-based sentiment analysis tool that is specifically made for social media text, we investigate the correlation between sentiment and subjectivity, and classes. Figure 8 shows a distribution and density of sentiment and subjectivity scores. Note that the two classes are laid on top of each other in both figures. A negative sentiment value naturally corresponds to negative sentiment, positive values correspond to positive sentiment and 0 corresponds to a neutral or non-sentimental tweet. A subjectivity score of 0 corresponds to an objective text, where 1 corresponds to maximum subjectivity.




Fig. 8: Distribution and sentiment

Observe that disaster tweets are more frequently objective and neutral. Additionally, non-disaster tweets carry more positive sentiment while disaster tweets are negative more often. Non-disaster tweets are also slightly skewed towards being more subjective than real disaster tweets. These observations all follow the intuition of disaster tweets being either mostly neutral and objective (reported news) or carry a negative sentiment.

Class	Negative	Neutral	Positive	Subjective
Non-Disaster	0.132	0.766	0.102	0.324
Disaster	0.173	0.777	0.049	0.265

TABLE II: Mean values of negative, neutral and positive sentiment as well as mean subjectivity of tweets for each class.

Table II shows the mean values of sentiment and subjectivity for the two classes. These values further quantifies the differences. Disaster tweets are 20% more objective, 71% less positive and 27% more negative than non-disaster tweets.

V. FEATURE EXPLORATION

VI. CONCLUSION

REFERENCES

- [1] T. Mikolov, I. Sutskever, K. Chen, G. Corrado, and J. Dean, "Distributed representations of words and phrases and their compositionality," *Advances in Neural Information Processing Systems*, vol. 26, 10 2013.

²<https://github.com/hermitdave/FrequencyWords>

³<https://github.com/cjhutto/vaderSentiment>

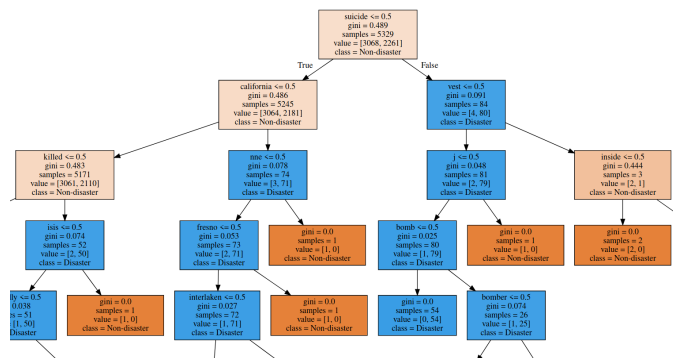


Fig. 9: First few nodes of a decision tree achieving an F1 score of 0.76.

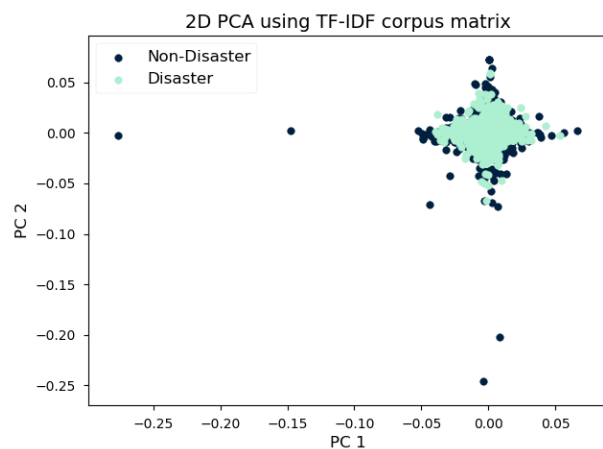


Fig. 10: 2 dimensional PCA using a TF-IDF corpus matrix.

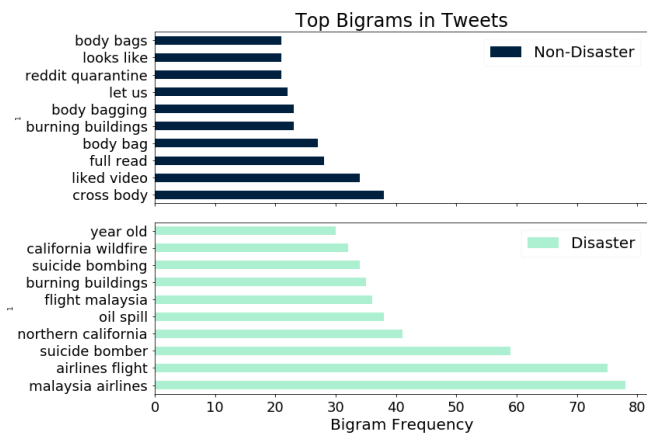


Fig. 11: Top 10 Bigrams.